

PLSC 503: “Multivariate Analysis for Political Research”

Exercise Eight

March 24, 2025

Part I

The *exponential distribution* is a simple one-parameter statistical distribution with PDF:

$$\Pr(X = x) = \lambda \exp(-\lambda x)$$

where $x \in [0, \infty)$ and $\lambda > 0$. You can read more about it [here](#). Using the exponential distribution as an example, your assignment¹ is to address:

1. *Consistency*: Using a Newton-Raphson optimization method, show via simulation that $\hat{\lambda}_{MLE} \rightarrow \lambda$ as $N \rightarrow \infty$.
2. *Invariance to reparameterization*: Reparameterize the exponential such that $\phi = 1/\lambda$, and show via simulation that $\hat{\phi}_{MLE} \rightarrow \phi$ as $N \rightarrow \infty$.
3. *Optimization*: Choose up to three alternative optimization algorithms, and discuss any differences in $\hat{\lambda}_{MLE}$ obtained by them. How do those differences change (if at all) as $N \rightarrow \infty$?

Part II

In the latter part of this exercise, we’ll compare parameter estimates from a simple, two-variable linear regression model obtained via OLS to those estimated through maximum likelihood. The data for this exercise (available as a `.csv` file on the course github site) are state-level data (excluding the District of Columbia) from 2009 ($N = 50$). They contain four variables:

- `State` is – unsurprisingly – the name of the state,
- `Obesity` is the percentage of adults in the state with a body mass index (BMI) greater than 30,
- `Temp` is the mean annual temperature for the state,
- `Beer` is the average per capita beer consumption in the state, in gallons.

The straightforward idea is to estimate the marginal associations of temperature and beer consumption with obesity; that is, treat `Obesity` as your “dependent variable.” In particular:

¹Hints: (a) you can use `rexp` to make random draws from an exponential distribution; (b) I am *not* asking you to write your own optimizer; you should use `maxLik` or `optim` or one of R’s other built-in optimizers.

1. Begin with the linear model:

$$\text{Obesity}_i = \beta_0 + \beta_1 \text{Temp}_i + \beta_2 \text{Beer}_i + u_i \quad (1)$$

and assume $u \sim N(0, \sigma^2)$. Write a short program/function to estimate $\beta_0, \beta_1, \beta_2$, and σ^2 via maximum likelihood.² Discuss your substantive “findings” to the best of your ability.

2. Estimate the same model via OLS, and compare the OLS and MLE results. What can you say about the various coefficient estimates obtained via the two methods? Their standard errors? The variance parameter σ^2 ? What can you say about the relative “goodness” of the two estimates?
3. Evaluate the following hypotheses for both the OLS- and MLE-estimated models:
 - (a) $\beta_1 = 0.3$
 - (b) $\beta_2 = \beta_1$
 - (c) $\sigma^2 = 10$

and discuss the implications of those evaluations.

4. Reestimate your MLE results, choosing a different optimization routine; so, if you used Newton-Raphson the first time around, switch to BHHH, or DFP, etc. Briefly discuss the differences (if any) between these “new” MLEs and those reported above.
5. Finally, reestimate your model under the assumption that the stochastic component follows a logistic distribution (that is, $u \sim \text{logistic}(0, s^2)$). You can read more about the logistic distribution in Evans et al., or [here](#). Report those estimates as well.

This assignment is due (in the usual electronic fashion) at 11:59 p.m. ET on Wednesday, April 2, 2025, and is worth 50 possible points.

²Note that, once again, you should feel free to use one of R’s built-in optimizers, although you’re also welcome to program your own if you’re so inclined.